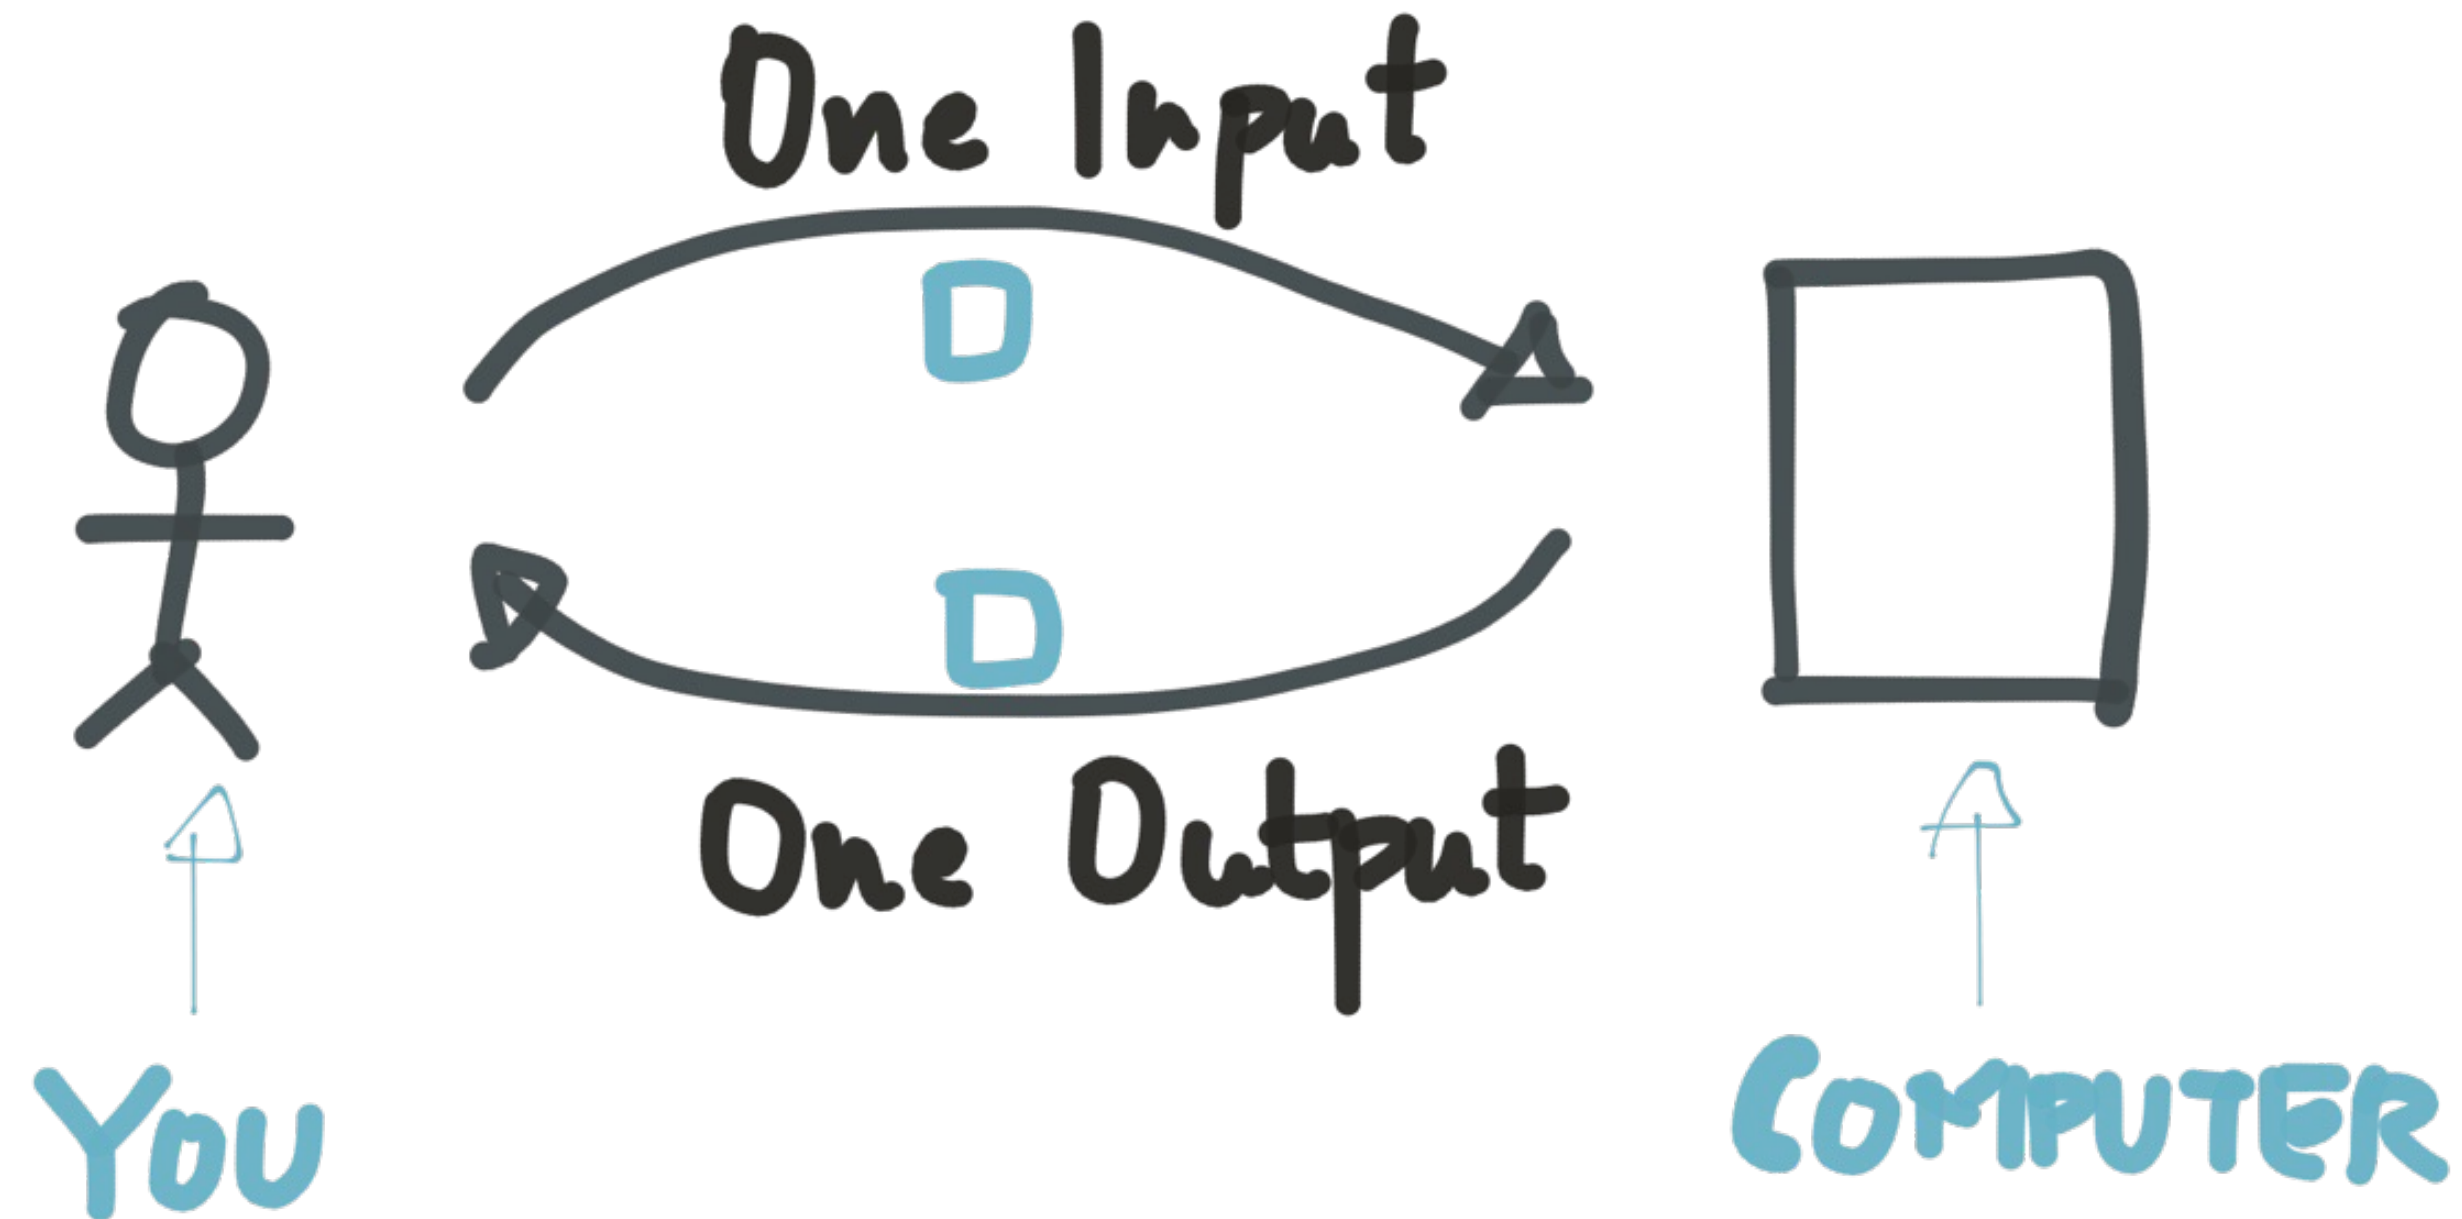



Kafka Streams

Stream processing Made Simple with Kafka

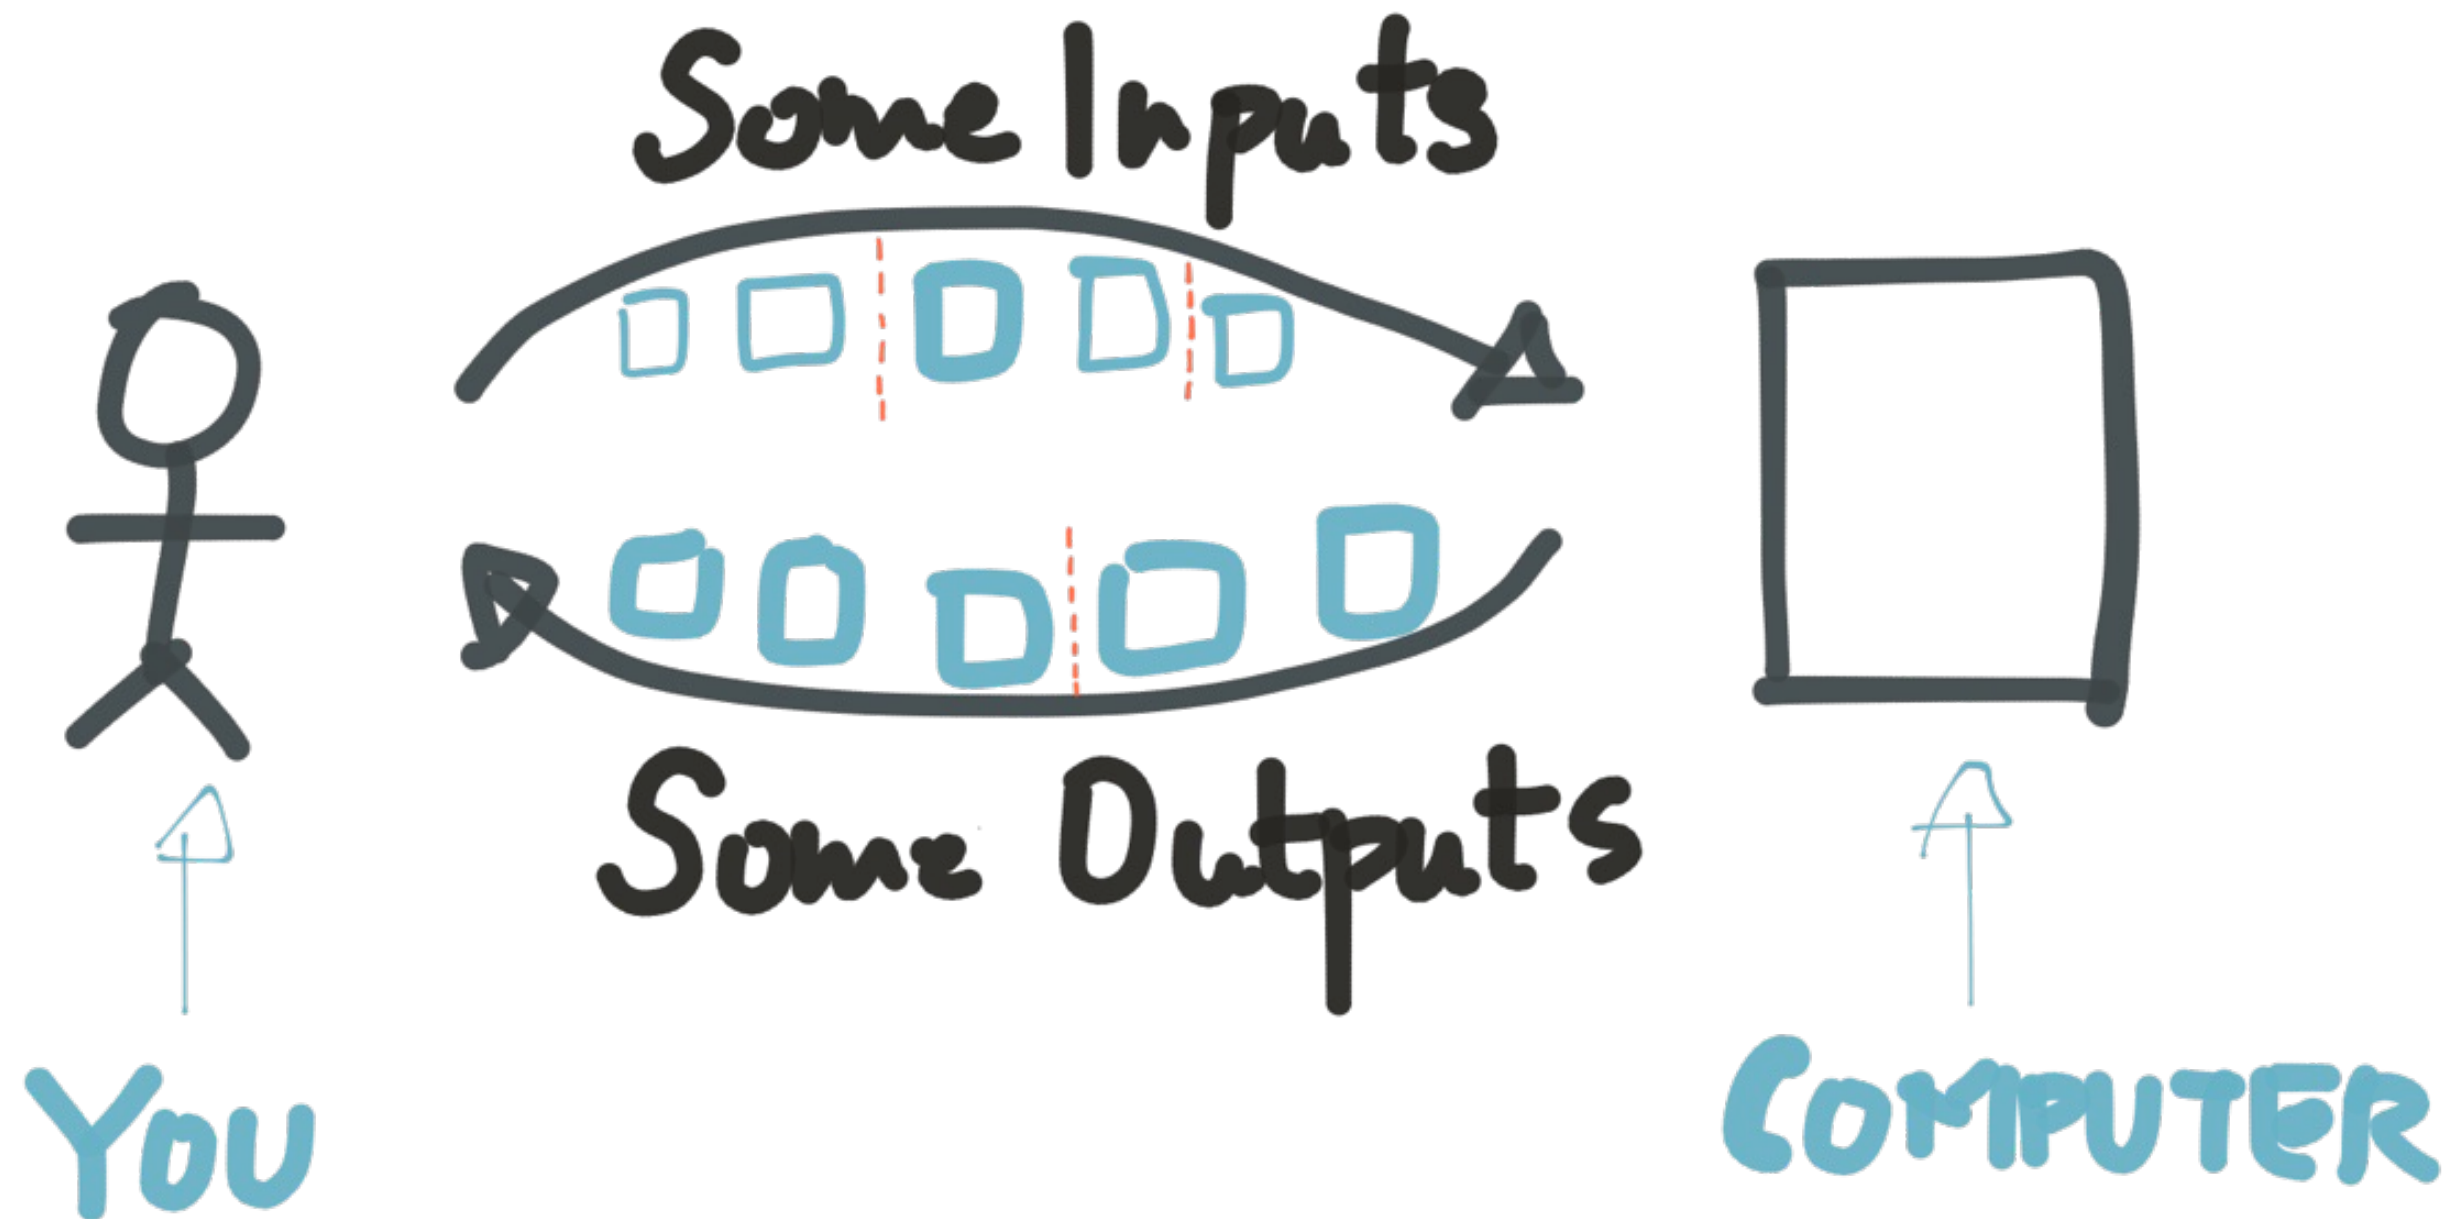
REQUEST / RESPONSE



BATCH



STREAM PROCESSING




Stream Processing

- *A different programming paradigm*
- *.. that brings computation to **unbounded** data*
- *.. with tradeoffs between **latency** / **cost** / **correctness***

Kafka Streams is a lightweight, yet powerful Java library for enriching, transforming, and processing real-time streams of data.

Kafka: Real-time Platforms

- *Persistent Buffering*
- *Logical Ordering*
- *Scalable “source-of-truth”*



APIs for moving data to and from Kafka


API	Topic interaction	Examples
Producer API	<i>Writing</i> messages to Kafka topics.	<ul style="list-style-type: none">• Filebeat• rsyslog• Custom producers
Consumer API	<i>Reading</i> messages from Kafka topics.	<ul style="list-style-type: none">• Logstash• kafkacat• Custom consumers
Connect API	<i>Connecting</i> external data stores, APIs, and filesystems to Kafka topics. Involves both <i>reading</i> from topics (sink connectors) and <i>writing</i> to topics (source connectors).	<ul style="list-style-type: none">• JDBC source connector• Elasticsearch sink connector• Custom connectors

Stream Processing with Kafka

- *Option 1: Do It Yourself !*

```
while (isRunning) {  
    // read some messages from Kafka  
    inputMessages = consumer.poll();  
  
    // do some processing...  
  
    // send output messages back to Kafka  
    producer.send(outputMessages);  
}
```


DIY STREAM PROCESSING



DIY Stream Processing is *Hard*

- *Ordering*
- *State Management*
- *Partitioning & Scalability*
- *Time, Window & Out-of-order Data*
- *Fault tolerance*
- *Re-processing*



Kafka Streams

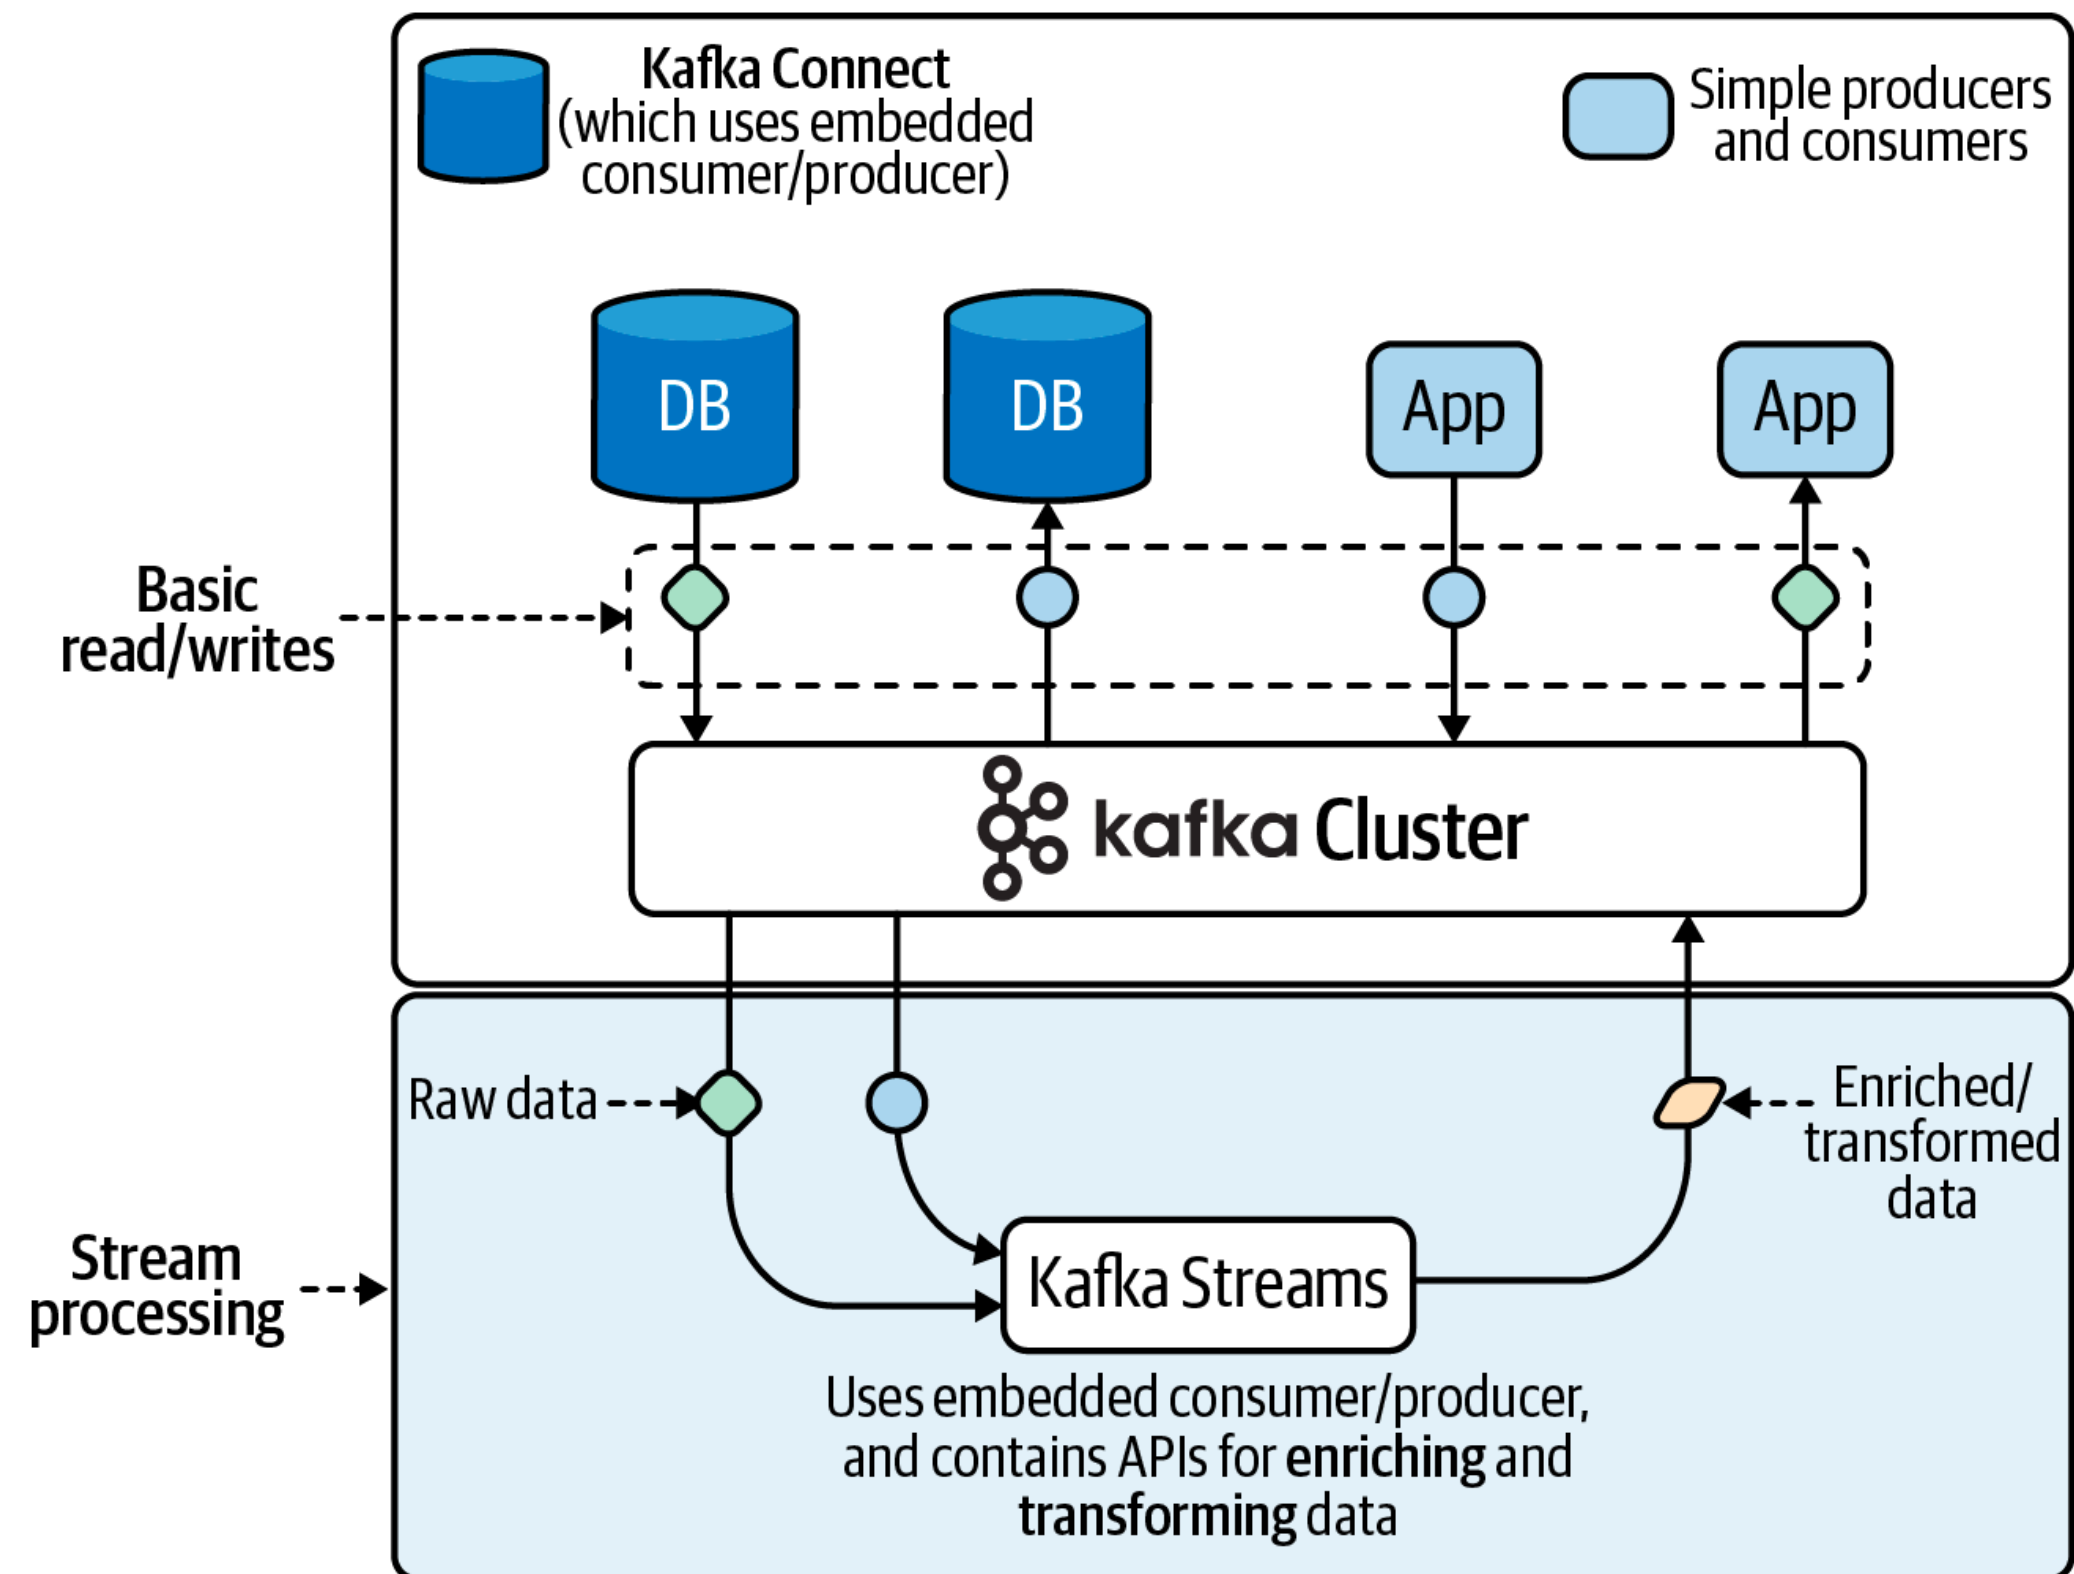


- *In Apache Kafka since v0.10, May 2016*
- *Powerful yet **easy-to-use** stream processing library*
 - *Event-at-a-time, Stateful*
 - *Windowing with out-of-order handling*
 - *Highly scalable, distributed, fault tolerant*
 - *and more..*

Anywhere, anytime

```
<dependency>  
  <groupId>org.apache.kafka</groupId>  
  <artifactId>kafka-streams</artifactId>  
  <version>2.10.0.0</version>  
</dependency>
```

Anywhere, anytime



Kafka Streams DSL

```
public static void main(String[] args) {  
    // specify the processing topology by first reading in a stream from a topic  
    KStream<String, String> words = builder.stream("topic1");  
  
    // count the words in this stream as an aggregated table  
    KTable<String, Long> counts = words.countByKey("Counts");  
  
    // write the result table to a new topic  
    counts.to("topic2");  
  
    // create a stream processing instance and start running it  
    KafkaStreams streams = new KafkaStreams(builder, config);  
    streams.start();  
}
```

Native Kafka Integration

```
Property cfg = new Properties();

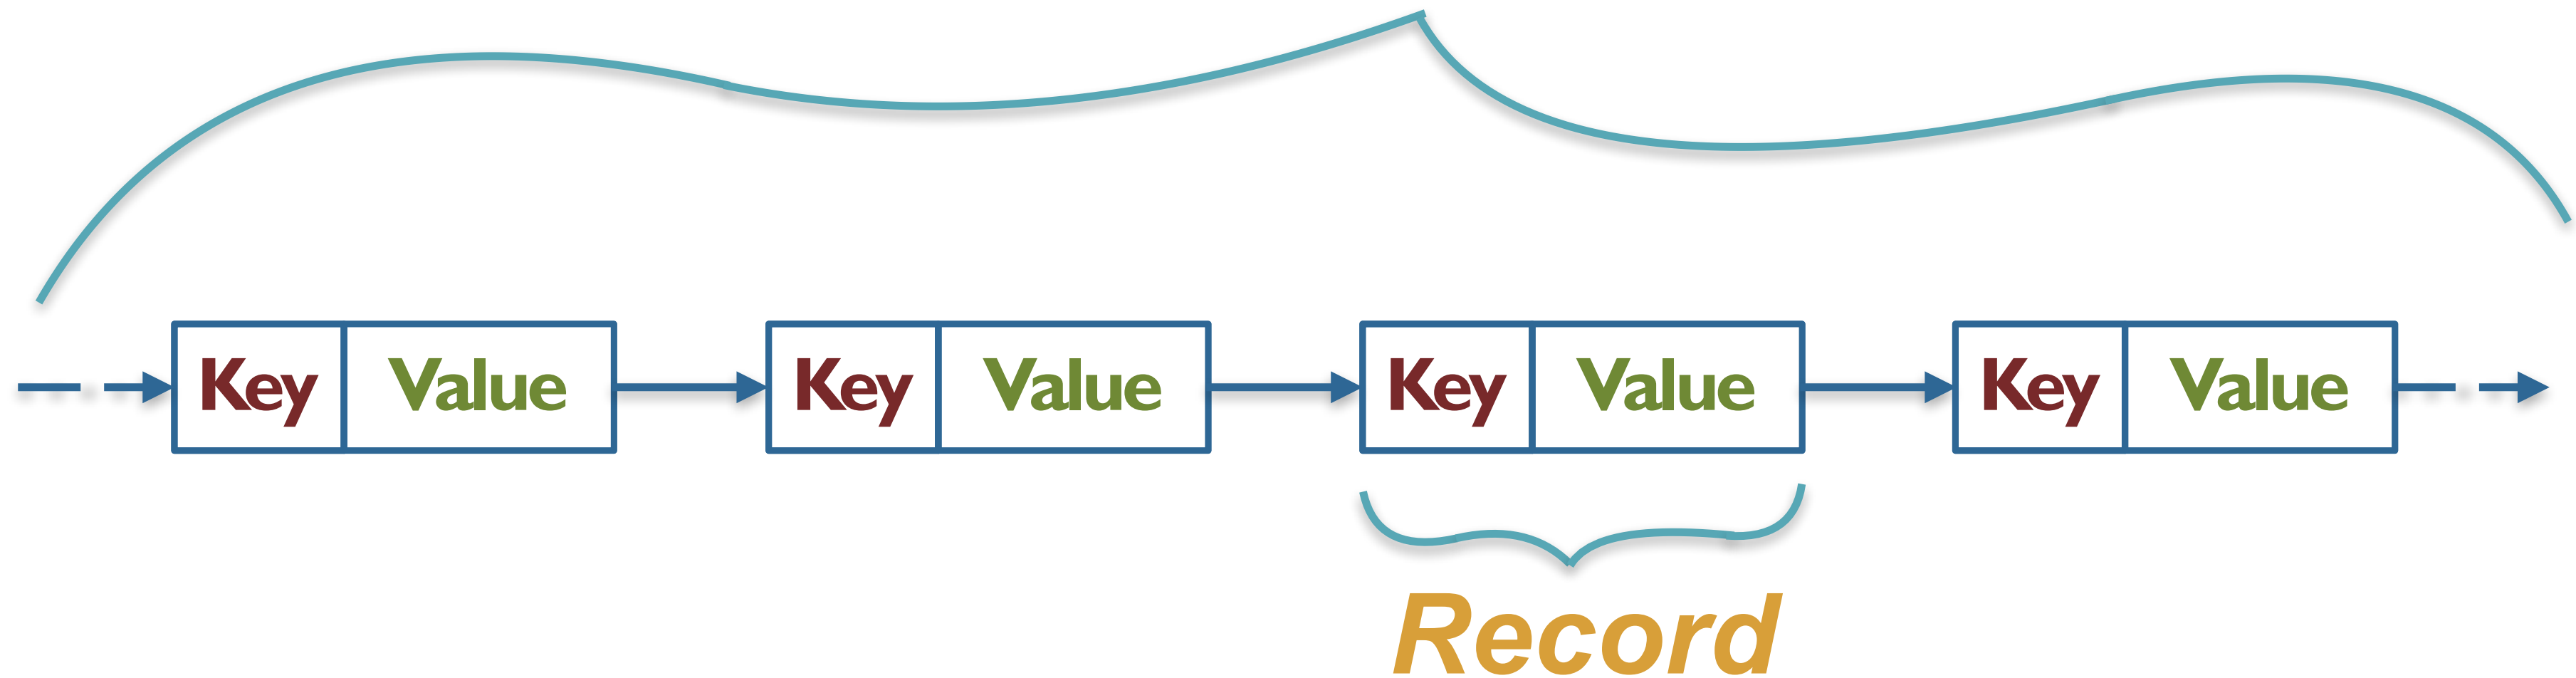
cfg.put(StreamsConfig.APPLICATION_ID_CONFIG, "my-streams-app");
cfg.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "broker1:9092");
cfg.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");
cfg.put(CommonClientConfigs.SECURITY_PROTOCOL_CONFIG, "SASL_SSL");
cfg.put(KafkaAvroSerDeConfig.SCHEMA_REGISTRY_URL_CONFIG, "registry:8081");

StreamsConfig config = new StreamsConfig(cfg);
final StreamsBuilder builder = new StreamsBuilder();
KafkaStreams streams = new KafkaStreams(builder, config);
```

Kafka Streams: **Key Concepts**

Stream and Records

Stream



Processor Topology

Source Processor

```
KStream<..> stream1 = builder.stream(  

```


```
KStream<..> stream2 = builder.stream(  

```

Sink Processor

```
aggregated.to(  

```



Stream API features:

- A high-level DSL that looks and feels like Java's streaming API
- A low-level Processor API that gives developers fine-grained control when they need it.
- Convenient abstractions for modeling data as either streams or tables
- The ability to join streams and tables, which is useful for data transformation and enrichment.
- Support for time-based operations, including windowing and periodic functions.
- Easy installation. It's just a library, so you can add Kafka Streams to any Java application.
- Scalability, reliability, maintainability

Uses Cases

- Financial data processing (Flipkart), purchase monitoring, fraud detection
- Algorithmic trading
- Stock market/crypto exchange monitoring
- Real-time inventory tracking and replenishment (Walmart)
- Event booking, seat selection (Ticketmaster)
- Email delivery tracking and monitoring (Mailchimp)

Lab : Using kafka Stream – 90 Minutes