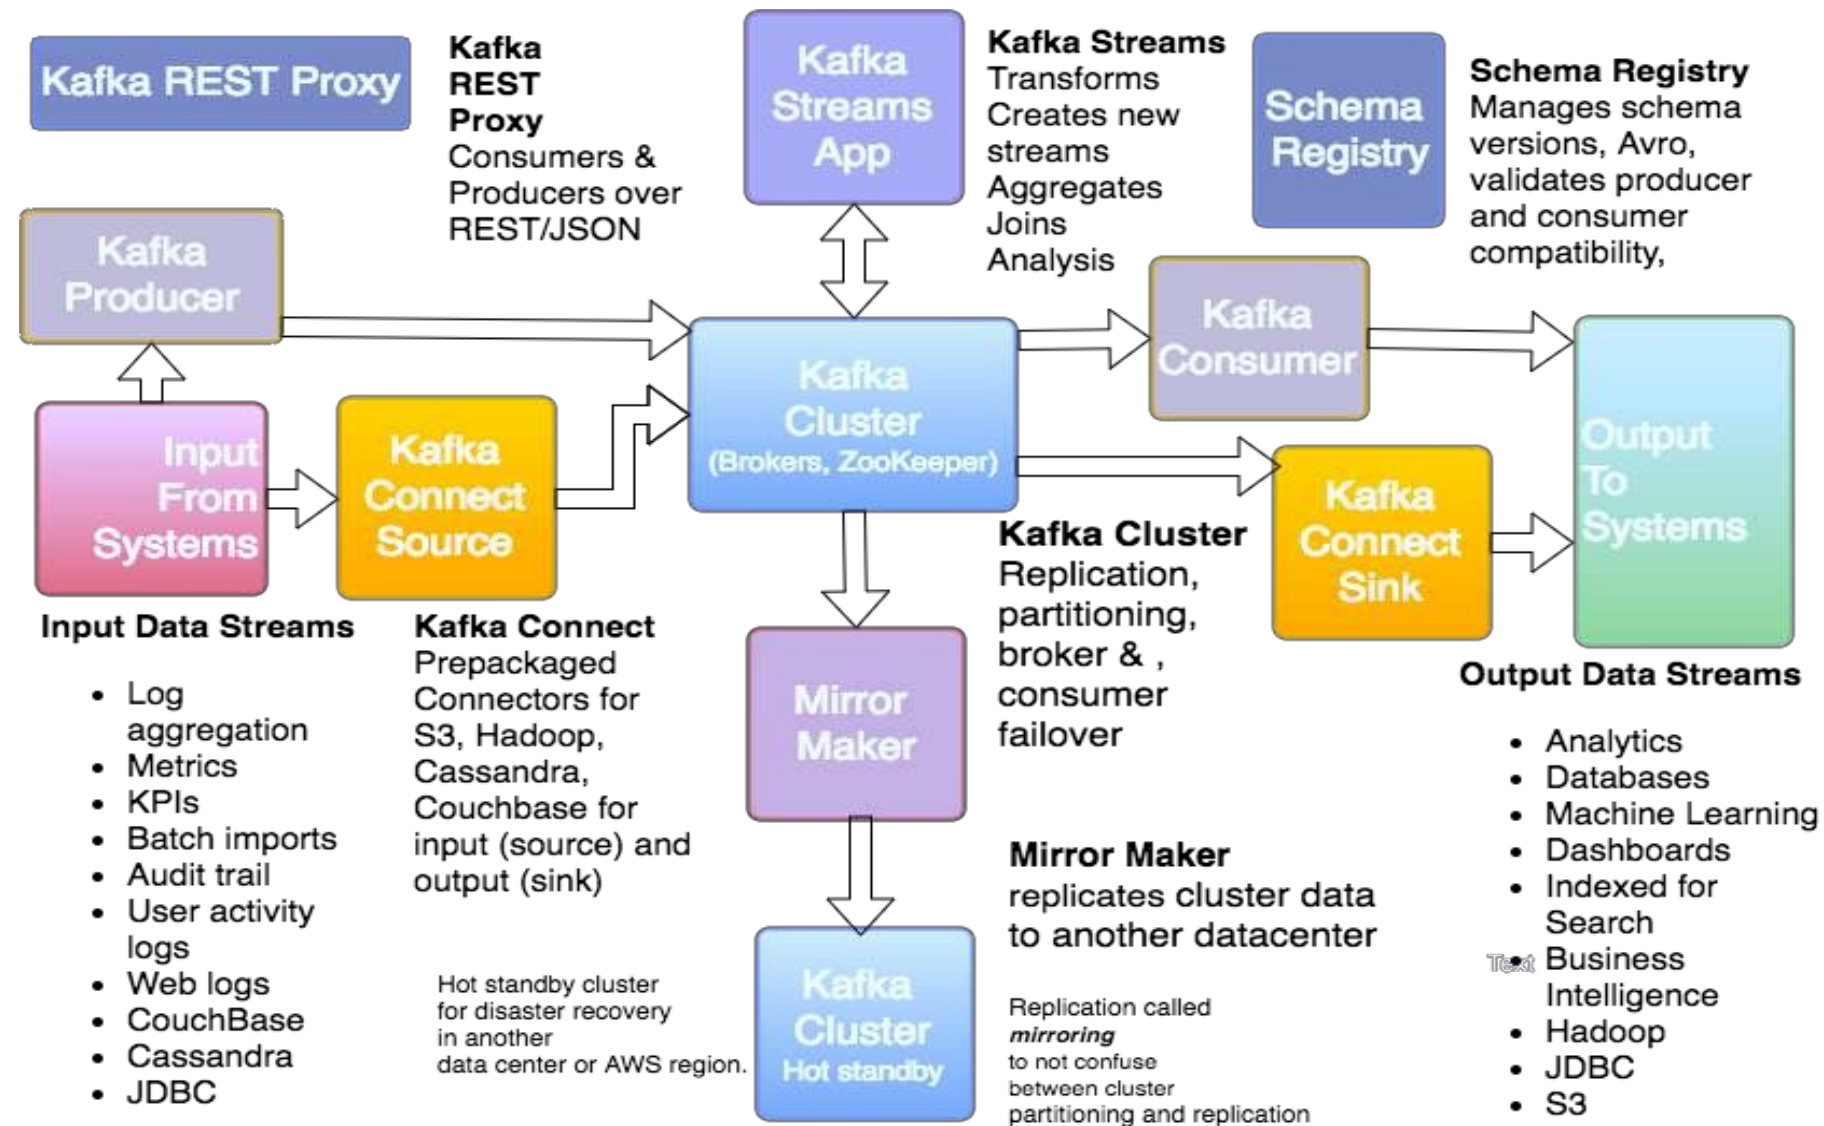
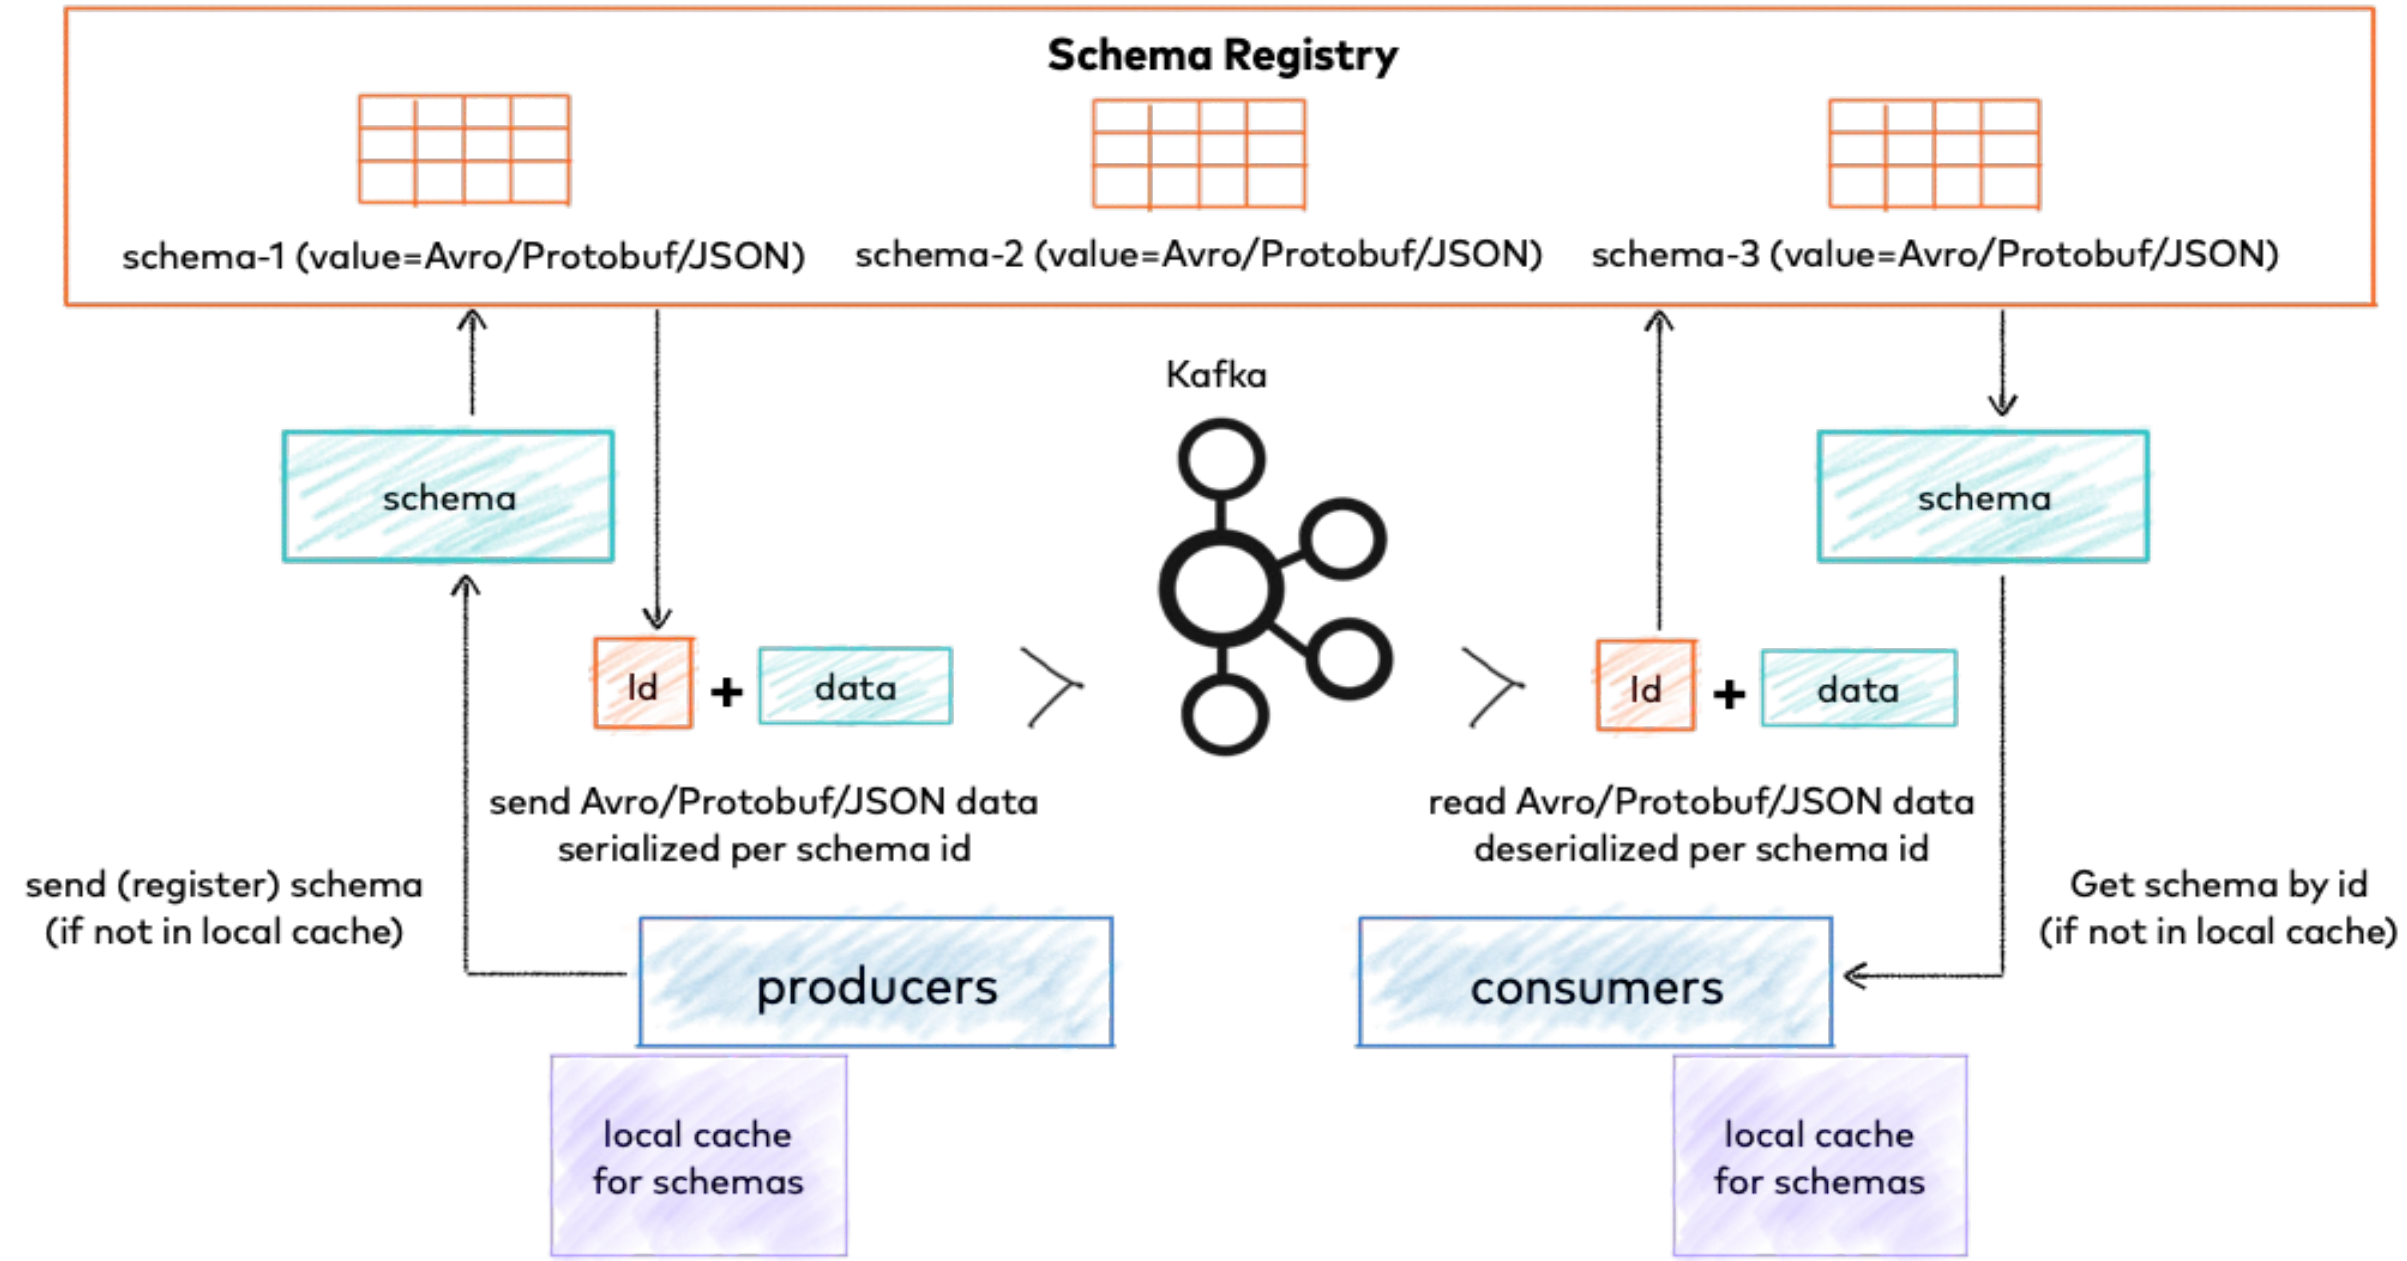


Kafka Ecosystem



Confluent Schema Registry



Schema Registry

The schema registry is a service that records the various schemas and their different versions as they evolve.

Producer and consumer clients retrieve schemas from the schema registry via HTTPS, store them locally in cache, and use them to serialize and deserialize messages sent to and received from Kafka.

This schema retrieval occurs only once for a given schema and from that point on the cached copy is relied upon.

Schema file - Avro

```
{
  "type": "record",
  "namespace": "io.confluent.developer.avro",
  "name": "Purchase",
  "fields": [
    {"name": "item", "type": "string"},
    {"name": "total_cost", "type": "double"},
    {"name": "customer_id", "type": "string"}
  ]
}
```

- The ***namespace*** field serves the same purpose as in the Protobuf file, preventing name collisions. The namespace also becomes the package name for the generated Java file.
- Fields are declared in a ***JSON*** array.

Confluent Schema Registry - Single Primary Architecture

ALL TOPICS >

transactions

Overview Messages Schema Configuration

Value Key

 Edit schema

 Version history

 Download

Format: AVRO Version: 1

```
1  {
2    "fields": [
3      {
4        "name": "id",
5        "type": "string"
6      },
7      {
8        "name": "amount",
9        "type": "double"
10     }
11   ],
12   "name": "Payment",
13   "namespace": "io.confluent.examples.clients.basicavro",
14   "type": "record"
15 }
```

Take note of the version, schema ID, and the fields.

Kafka Connect : Source & Sink Connectors

Tos

